



# Experiences of Porting Structured and Unstructured Stencil Applications to FPGA using SYCL

Zadok Storkey, Steven Wright, Ian Gray

# 1. Context

1. Context
2. Approach
3. Results & Discussion
4. Future Work

# Context

**The potential benefits to reducing power consumption of HPC systems are clear:**

- Environmental - i.e. smaller carbon footprint
- Financial - i.e. lower power and cooling cost
- Infrastructural - i.e. less burden on energy infrastructure and water supply

**Finding alternative low-energy accelerators is a good way of achieving this aim.**

# Context

FPGAs potentially fit this niche:

- Low energy
- Highly parallel

However, there are significant challenges to effectively utilise them:

- How can we use them *portably*?
- How can we use them *productively*?



## **2. Approach**

- 1. Context**
- 2. Approach**
- 3. Results & Discussion**
- 4. Future Work**

# Approach


## Outline

- Goal: Evaluate the process and results of adapting an existing CPU application portably for FPGA
- Application: Structured and Unstructured version of the *heat* mini application [1]
- Programming model: Intel OneAPI implementation of SYCL
- Targets: AMD EPYC 7713 CPU, Intel Stratix 10 MX FPGA
- Environment: EPCC Testbed

# Approach

## Application

- Based on the *heat* [1] mini-application
- Solves the heat equation over a 2D grid, using a finite difference method
- Uses an iterative approach with a 5-point stencil
- Two versions: structured grid and unstructured grid
- Unstructured version further divided by degree of memory shuffle



# Approach

**Implemented in three ways:**

- CPU (Parallel For) – SYCL controls loop flow and assigns different iterations to different CPU cores
- FPGA (Parallel For) – SYCL controls loop flow and runs the kernel, which is synthesised and implemented onto an FPGA
- FPGA (Single Task) – The kernel maintains control of the loop flow, and is synthesised and implemented onto an FPGA

# 3. Results & Discussion

1. Context
2. Approach
- 3. Results & Discussion**
4. Future Work


# Results & Discussion

## Discussion – How single source is the application actually?

```
1 cgh.parallel_for<class solve_kernel>(sycl::range<2>(nrows, ncols), [=](sycl::id
2     <2> idx) [[sycl::reqd_work_group_size(16, 16)]]) {
3     const long row = idx[0];
4     const long col = idx[1];
5     const double temp_centre = get_previous_temps[row][col];
6     double temp_adj_sum = 0;
7     for (int face = 0; face < 4; face++) {
8         const long adj_row = get_adj_row[row][col][face];
9         const long adj_column = get_adj_col[row][col][face];
10        if (adj_row != Neighbour::BOUNDARY) {
11            adj_temperature_sum += get_previous_temps[adj_row][adj_col];
12        }
13    }
14    set_next_temps[row][column] = r2 * temp_centre + r * temp_adj_sum;
15});
```

Table 3: Code divergence of the kernel by target in lines  
**(structured | unstructured)**

	CPU	FPGA (parallel-for)	FPGA (single-task)
CPU	-	1   3	6   8
FPGA (parallel-for)	1   3	-	6   6
FPGA (single-task)	6   8	6   6	-



# Results & Discussion

## Discussion – How single source is the application actually?

```
1 cgh.parallel_for<class solve_kernel>(sycl::range<2>(nrows, ncols), [=](sycl::id
2     <2> idx) [[sycl::reqd_work_group_size(16, 16)]]
3 {
4     const long row = idx[0];
5     const long col = idx[1];
6     const double temp_centre = get_previous_temps[row][col];
7     double temp_adj_sum = 0;
8     for (int face = 0; face < 4; face++) {
9         const long adj_row = get_adj_row[row][col][face];
10        const long adj_column = get_adj_col[row][col][face];
11        if (adj_row != Neighbour::BOUNDARY) {
12            adj_temperature_sum += get_previous_temps[adj_row][adj_col];
13        }
14    }
15    set_next_temps[row][column] = r2 * temp_centre + r * temp_adj_sum;
16});
```

```
1 cgh.parallel_for<class solve_kernel>(sycl::range<2>(nrows, ncols), [=](sycl::id
2     <2> idx)
3     [[intel::kernel_args_restrict, intel::num_simd_work_items(16), sycl::reqd
4     _work_group_size(1, 16)]]
5 {
6     const long row = idx[0];
7     const long column = idx[1];
8     const double temp_centre = get_previous_temps[row][col];
9     double temp_adj_sum = 0;
10    #pragma unroll
11    #if intel::lvdep
12    for (int face = 0; face < 4; face++) {
13        const long adj_row = get_adj_row[row][col][face];
14        const long adj_column = get_adj_col[row][col][face];
15        if (adj_row != Neighbour::BOUNDARY) {
16            adj_temperature_sum += get_previous_temps[adj_row][adj_col];
17        }
18    }
19    set_next_temps[row][column] = r2 * temp_centre + r * temp_adj_sum;
20});
```

Table 3: Code divergence of the kernel by target in lines  
**(structured | unstructured)**

	CPU	FPGA (parallel-for)	FPGA (single-task)
CPU	-	1   3	6   8
FPGA (parallel-for)	1   3	-	6   6
FPGA (single-task)	6   8	6   6	-

# Results & Discussion

## Discussion – How single source is the application actually?

```
1 cgh.parallel_for<class solve_kernel>(sycl::range<2>(nrows, ncols), [=](sycl::id
2   <2> idx) [[sycl::reqd_work_group_size(16, 16)]]) {
3   const long row = idx[0];
4   const long col = idx[1];
5   const double temp_centre = get_previous_temps[row][col];
6   double temp_adj_sum = 0;
7   for (int face = 0; face < 4; face++) {
8     const long adj_row = get_adj_row[row][col][face];
9     const long adj_column = get_adj_col[row][col][face];
10    if (adj_row != Neighbour::BOUNDARY) {
11      adj_temperature_sum += get_previous_temps[adj_row][adj_col];
12    }
13  }
14  set_next_temps[row][column] = r2 * temp_centre + r * temp_adj_sum;
15 });
16 }
```

```
1 cgh.parallel_for<class solve_kernel>(sycl::range<2>(nrows, ncols), [=](sycl::id
2   <2> idx) [[intel::kernel_args_restrict, intel::num_simd_work_items(16), sycl::reqd
3   _work_group_size(1, 16)]]) {
4   const long row = idx[0];
5   const long column = idx[1];
6   const double temp_centre = get_previous_temps[row][col];
7   double temp_adj_sum = 0;
8   #pragma unroll
9   #if intel::ivdep
10  for (int face = 0; face < 4; face++) {
11    const long adj_row = get_adj_row[row][col][face];
12    const long adj_column = get_adj_col[row][col][face];
13    if (adj_row != Neighbour::BOUNDARY) {
14      adj_temperature_sum += get_previous_temps[adj_row][adj_col];
15    }
16  }
17  set_next_temps[row][column] = r2 * temp_centre + r * temp_adj_sum;
18 });
19 }
```

Table 3: Code divergence of the kernel by target in lines  
**(structured | unstructured)**

	CPU	FPGA (parallel-for)	FPGA (single-task)
CPU	-	1   3	6   8
FPGA (parallel-for)	1   3	-	6   6
FPGA (single-task)	6   8	6   6	-

```
1 cgh.single_task<class solve_kernel>([=]() [[intel::kernel_args_restrict]] {
2   [[intel::ivdep]]
3   for (long row = 0; row < nrows; row++) {
4     #pragma unroll 32
5     [[intel::ivdep]]
6     for (long column = 0; column < ncols; column++) {
7       const double temp_centre = get_previous_temps[row][col];
8       double temp_adj_sum = 0;
9       #pragma unroll
10      [[intel::ivdep]]
11      for (int face = 0; face < 4; face++) {
12        const long adj_row = get_adj_row[row][col][face];
13        const long adj_column = get_adj_col[row][col][face];
14        if (adj_row != Neighbour::BOUNDARY) {
15          adj_temperature_sum += get_previous_temps[adj_row][adj_col];
16        }
17      }
18      set_next_temps[row][column] = r2 * temp_centre + r * temp_adj_sum;
19    });
20  });
21 }):
```

# Results & Discussion

## Discussion – How easy was it to develop with FPGAs as a target

- The features of the toolchain helped with the productivity issues caused by the time taken to build FPGA images:
  - Fast-compile
  - Image re-use
  - Pre-synthesis reports
- Nevertheless, build times remain a challenge

**Table 4: Build times (hh:mm:ss) by FPGA target variation and grid type**

	Parallel-for	Single-task
Structured	03:19:28	04:38:37
Unstructured	02:47:57	06:36:37

# Results & Discussion

## Results – Solve times

**Table 1: Solve times (s) by target and grid type**

	CPU	FPGA (parallel-for)	FPGA (single-task)
Structured	15.27	252.90	208.42
Unstructured (No Shuffle)	88.60	8566.28	2074.41
Unstructured (Shuffle ( $0.5 \cdot n_{cells}$ ))	91.93	9384.59	1790.39
Unstructured (Shuffle ( $1.0 \cdot n_{cells}$ ))	93.91	9578.58	1795.50
Unstructured (Shuffle ( $2.0 \cdot n_{cells}$ ))	96.27	9676.64	1795.82
Unstructured (Full Shuffle)	95.96	9882.52	2899.41

# Results & Discussion

## Results – Solve times

**Table 1: Solve times (s) by target and grid type**

	CPU	FPGA (parallel-for)	FPGA (single-task)
Structured	15.27	252.90	208.42
Unstructured (No Shuffle)	88.60	8566.28	2074.41
Unstructured (Shuffle ( $0.5 \cdot n_{cells}$ ))	91.93	9384.59	1790.39
Unstructured (Shuffle ( $1.0 \cdot n_{cells}$ ))	93.91	9578.58	1795.50
Unstructured (Shuffle ( $2.0 \cdot n_{cells}$ ))	96.27	9676.64	1795.82
Unstructured (Full Shuffle)	95.96	9882.52	2899.41

# Results & Discussion

## Results – Solve times

**Table 1: Solve times (s) by target and grid type**

	CPU	FPGA (parallel-for)	FPGA (single-task)
Structured	15.27	252.90	208.42
Unstructured (No Shuffle)	88.60	8566.28	2074.41
Unstructured (Shuffle ( $0.5 \cdot n_{cells}$ ))	91.93	9384.59	1790.39
Unstructured (Shuffle ( $1.0 \cdot n_{cells}$ ))	93.91	9578.58	1795.50
Unstructured (Shuffle ( $2.0 \cdot n_{cells}$ ))	96.27	9676.64	1795.82
Unstructured (Full Shuffle)	95.96	9882.52	2899.41

# Results & Discussion

## Results – Energy usage

**Table 2: Energy Usage (kJ) by target and grid type**

	CPU	FPGA (parallel-for)	FPGA (single-task)
Structured	2.92	6.02*	8.82*
Unstructured (No Shuffle)	17.33	212.70*	91.54*
Unstructured (Shuffle ( $0.5 \cdot n_{cells}$ ))	18.23	233.02*	79.01*
Unstructured (Shuffle ( $1.0 \cdot n_{cells}$ ))	18.64	237.84*	79.24*
Unstructured (Shuffle ( $2.0 \cdot n_{cells}$ ))	18.78	240.27*	79.25*
Unstructured (Full Shuffle)	24.52	245.38*	127.95*

\* estimates provided by Quartus

# Results & Discussion

## Results – Energy usage

**Table 2: Energy Usage (kJ) by target and grid type**

	CPU	FPGA (parallel-for)	FPGA (single-task)
Structured	2.92	6.02*	8.82*
Unstructured (No Shuffle)	17.33	212.70*	91.54*
Unstructured (Shuffle ( $0.5 \cdot n_{cells}$ ))	18.23	233.02*	79.01*
Unstructured (Shuffle ( $1.0 \cdot n_{cells}$ ))	18.64	237.84*	79.24*
Unstructured (Shuffle ( $2.0 \cdot n_{cells}$ ))	18.78	240.27*	79.25*
Unstructured (Full Shuffle)	24.52	245.38*	127.95*

\* estimates provided by Quartus

# Results & Discussion

## Results – Energy usage

**Table 2: Energy Usage (kJ) by target and grid type**

	CPU	FPGA (parallel-for)	FPGA (single-task)
Structured	2.92	6.02*	8.82*
Unstructured (No Shuffle)	17.33	212.70*	91.54*
Unstructured (Shuffle ( $0.5 \cdot n_{cells}$ ))	18.23	233.02*	79.01*
Unstructured (Shuffle ( $1.0 \cdot n_{cells}$ ))	18.64	237.84*	79.24*
Unstructured (Shuffle ( $2.0 \cdot n_{cells}$ ))	18.78	240.27*	79.25*
Unstructured (Full Shuffle)	24.52	245.38*	127.95*

\* estimates provided by Quartus

# Results & Discussion

## Discussion – Solve times and Energy Usage

- All FPGA builds perform rather poorly
- In terms of energy usage, FPGAs perform closer to CPUs, but still perform worse than expected
- Low computational intensity means that problem is likely limited more by bandwidth than compute
  - Supported by hardware resource utilisation amounts
- Age of hardware is also a consideration

## **4. Future Work**

- 1. Context**
- 2. Approach**
- 3. Results & Discussion**
- 4. Future Work**

# Future work

## Future work

- Expand Hardware Targets
  - Compare GPU
  - Compare more recent Intel FPGAs
  - Compare AMD FPGA using SYCL for Vitis
- Better energy data collection

# Future work

## Future work

- Investigate the causes behind the poor performance
  - Investigate behaviour of memory transfers during kernel execution
  - Experiment with increasing the arithmetic intensity of the kernel
- Expand to larger, more complex unstructured applications

# Experiences of Porting Structured and Unstructured Stencil Applications to FPGA using SYCL

Zadok Storkey, Steven Wright, Ian Gray